

DS HW4 Tree Based Models

JunLu

4/16/2019

Problem 1

a) Fit a regression tree

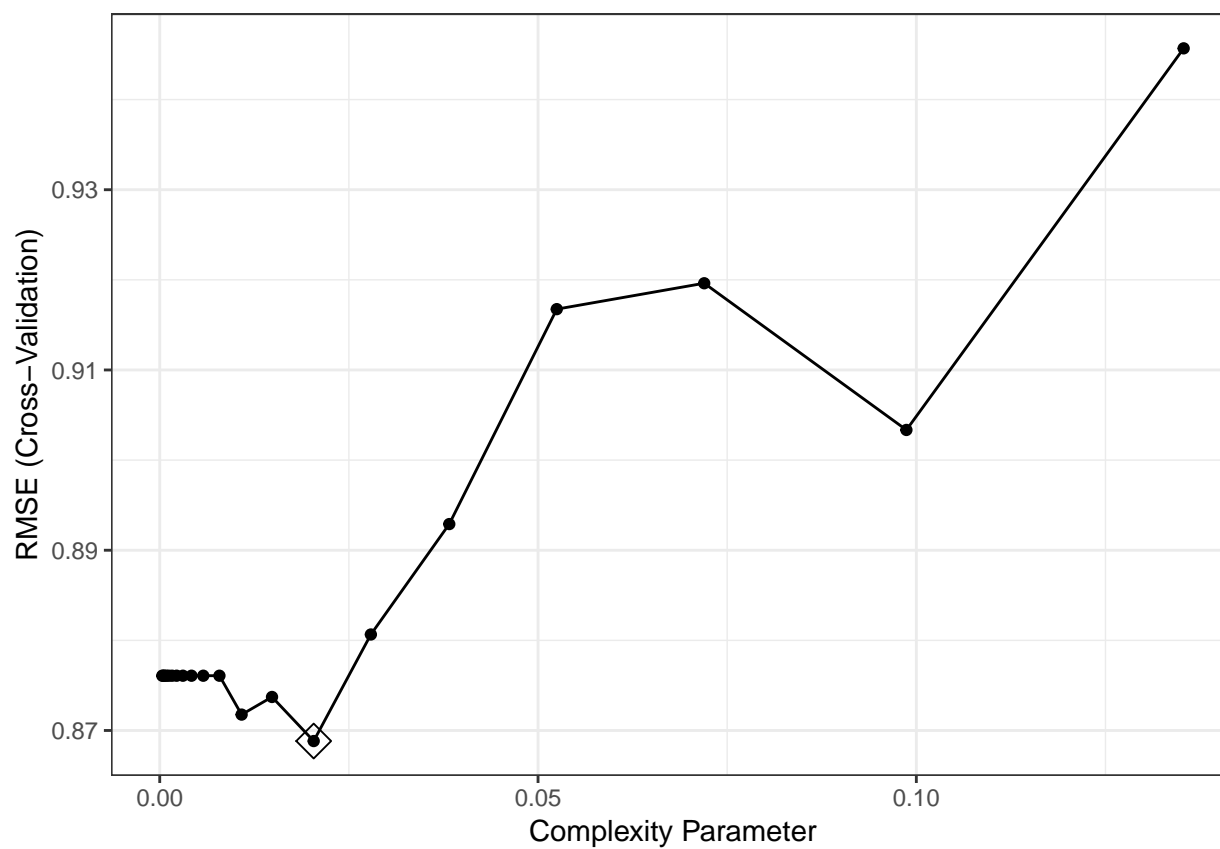
The lowest cross-validation error

```
ctrl = trainControl(method = "cv")
```

```
set.seed(1)
```

```
tree_fit = train(lpsa~.,  
  data = Prostate,  
  method = "rpart",  
  tuneGrid = data.frame(cp = exp(seq(-8,-2, length = 20))),  
  trControl = ctrl)
```

```
ggplot(tree_fit, highlight = TRUE)
```



```
tree_fit$bestTune
```

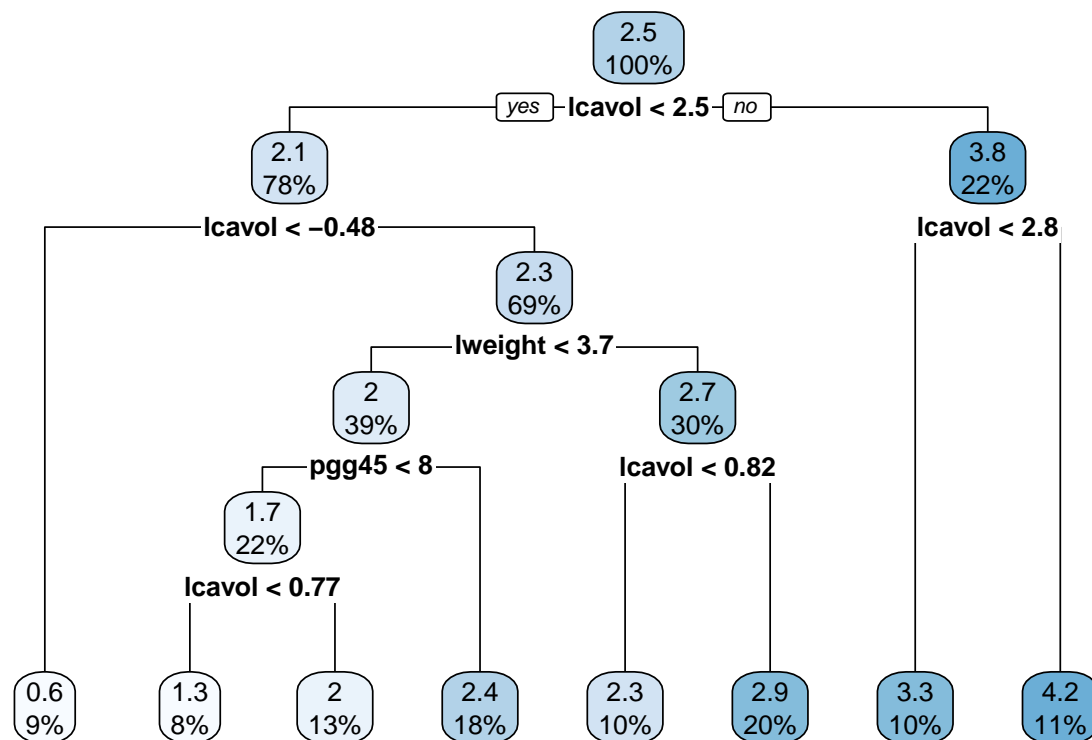
```
##          cp
```

```
## 14 0.02034873
```

```
tree_fit$finalModel$cptable
```

```
##          CP nsplit rel error
## 1 0.34710828    0 1.0000000
## 2 0.18464743    1 0.6528917
## 3 0.05931585    2 0.4682443
## 4 0.03475635    3 0.4089284
## 5 0.03460901    4 0.3741721
## 6 0.02156368    5 0.3395631
## 7 0.02146995    6 0.3179994
## 8 0.00000000    7 0.2965295
```

```
rpart.plot(tree_fit$finalModel)
```

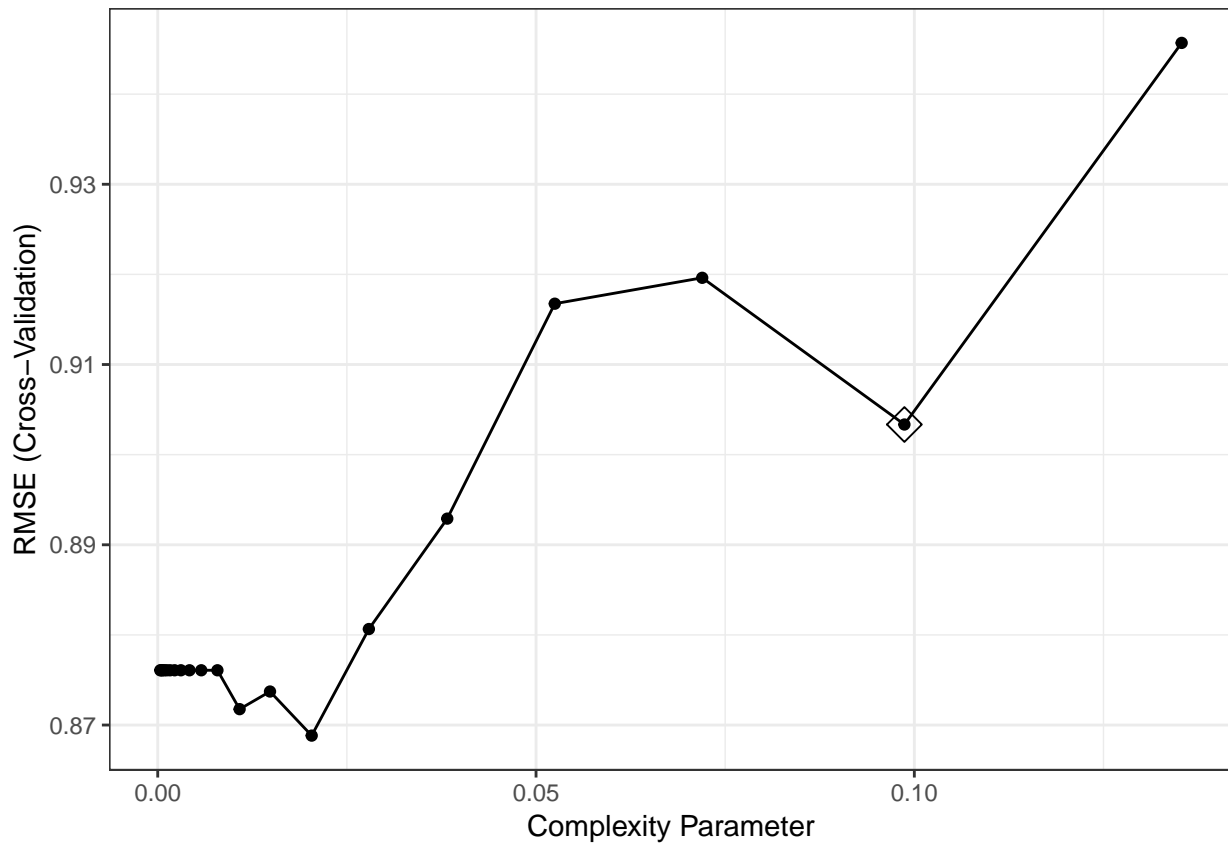


The tree size corresponds to the lowest cross-validation error is 8

The 1 SE rule

```
set.seed(1)
tree_fit_2 = train(lpsa~.,
  data = Prostate,
  method = "rpart",
  tuneGrid = data.frame(cp = exp(seq(-8, -2, length = 20))),
  trControl = trainControl(method = "cv",
    number = 10,
    selectionFunction = "oneSE"))

ggplot(tree_fit_2, highlight = TRUE)
```



```
tree_fit_2$finalModel$cptable
```

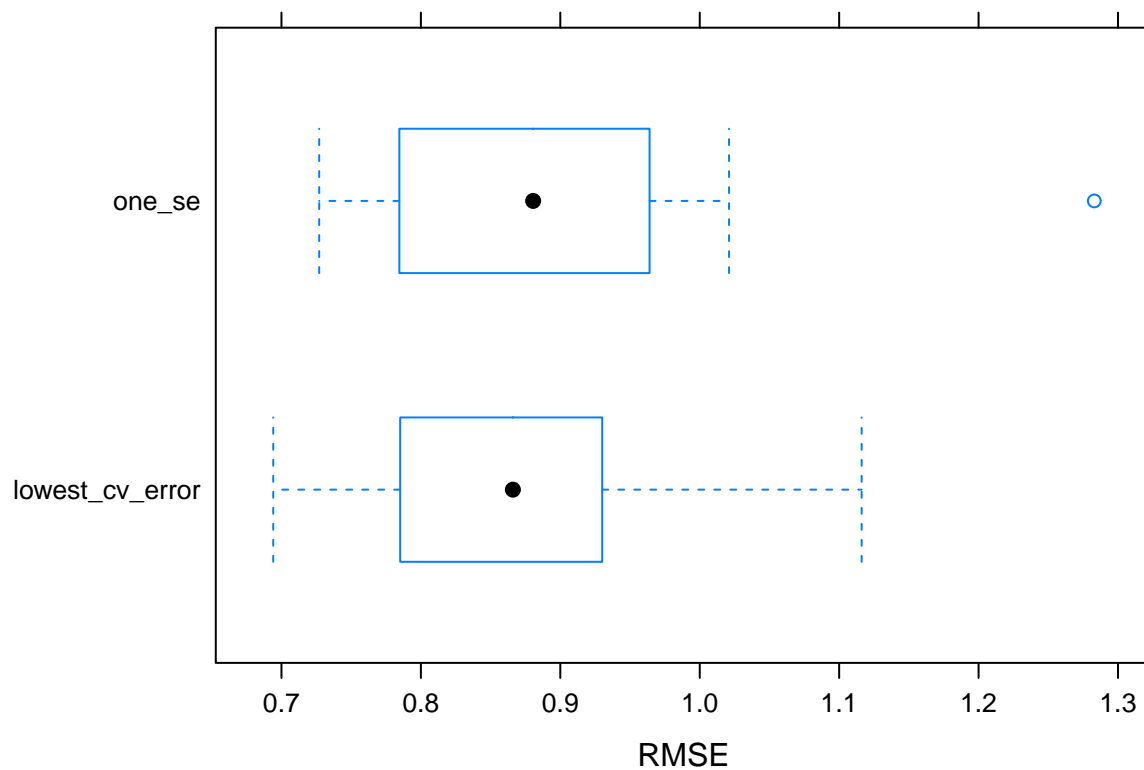
```
##          CP nsplit rel error
## 1 0.34710828      0 1.0000000
## 2 0.18464743      1 0.6528917
## 3 0.09868824      2 0.4682443
```

The tree size obtained using the 1 SE rule is 3.

The two tree sizes obtained by different selection functions are different.

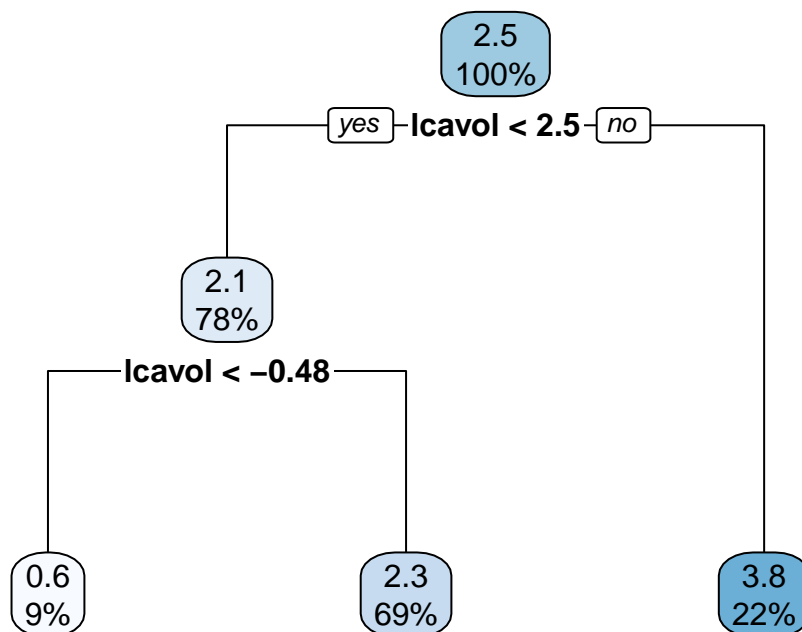
b) Choose one decision tree model

```
resamp <- resamples(list(lowest_cv_error = tree_fit,
                        one_se = tree_fit_2
                      ))
bwplot(resamp, metric = "RMSE")
```



I used tree size 3 as it has a similar performance with size 8 and it is simpler.

```
rpart.plot(tree_fit_2$finalModel)
```



The first terminal node in the plot: When the lcavol is smaller than -0.48 (firstly smaller than 2.5), the predicted value (or the mean of observations in this terminal node) is 0.6. This terminal node contains 9% training data observations.

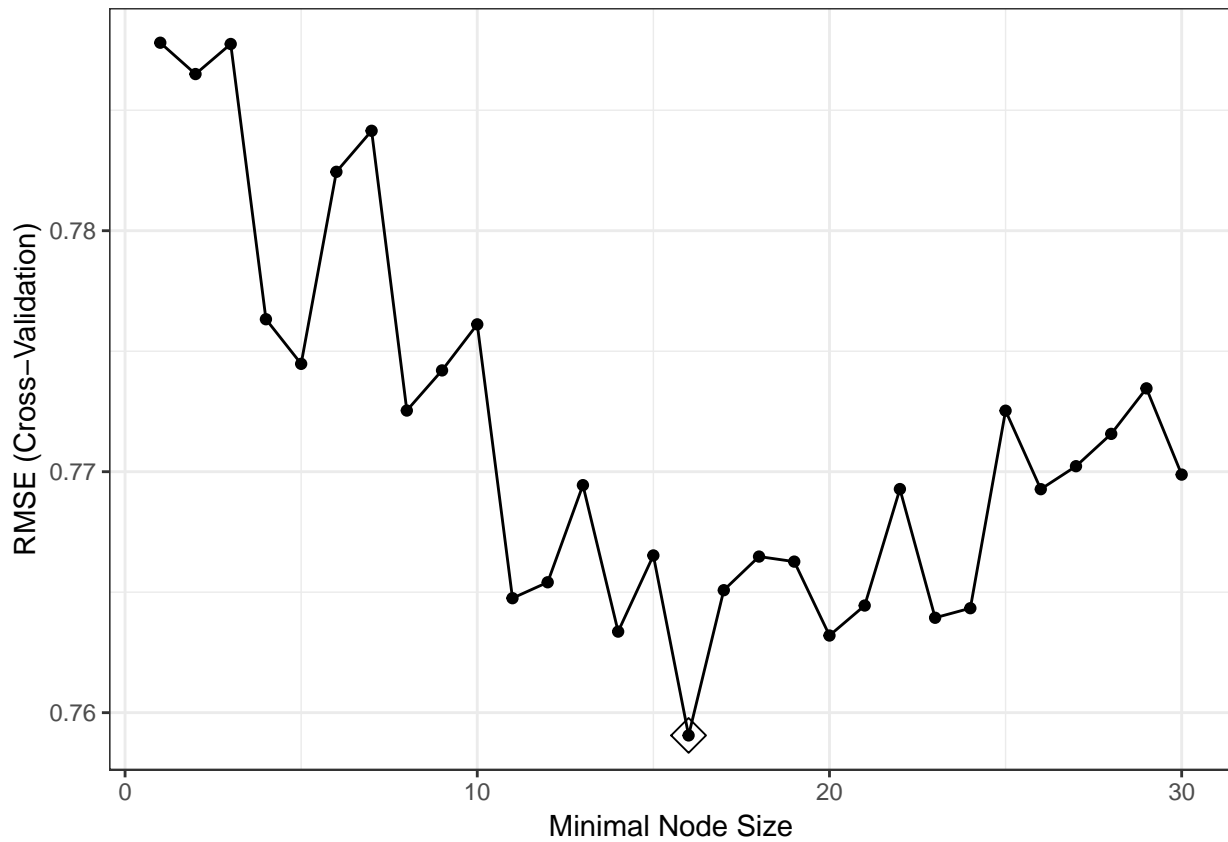
c) Bagging

```
set.seed(1)

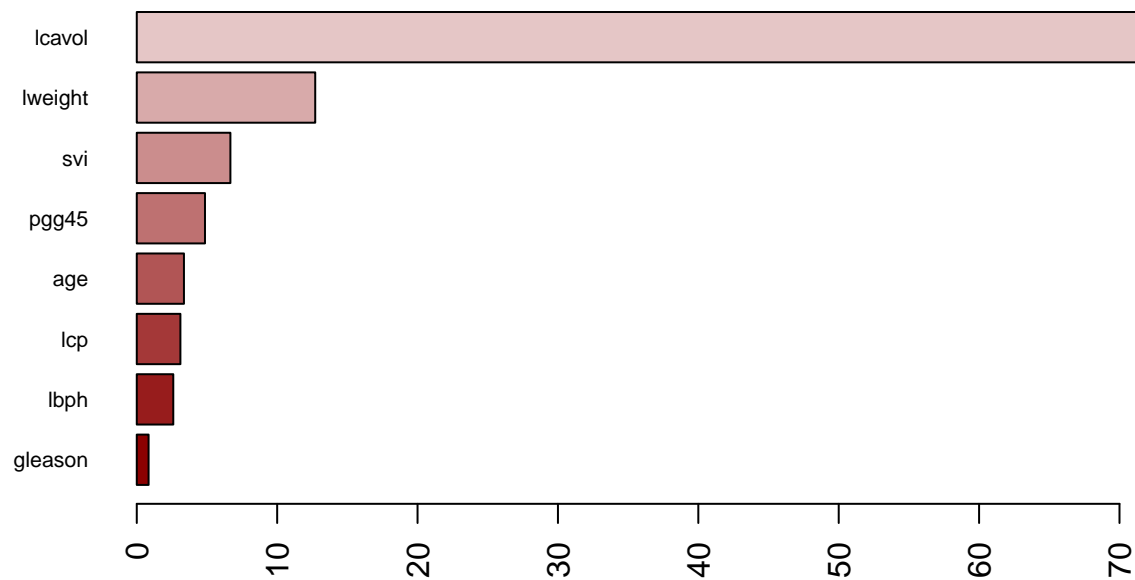
bagging.grid <- expand.grid(mtry = 8,
                           splitrule = "variance",
                           min.node.size = 1:30
                           )

bag_fit = train(lpsa~.,
                data = Prostate,
                method = "ranger",
                tuneGrid = bagging.grid,
                importance = "impurity",
                trControl = ctrl)

ggplot(bag_fit, highlight = T)
```



```
barplot(sort(ranger::importance(bag_fit$finalModel),
             decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(
          colors = c("darkred", "white", "darkblue"))(19))
```



The lcavol is the most importance variable in this bagging model.

Importance: lcavol > lweight > svi > pgg45 > age > lcp > lbph > gleason

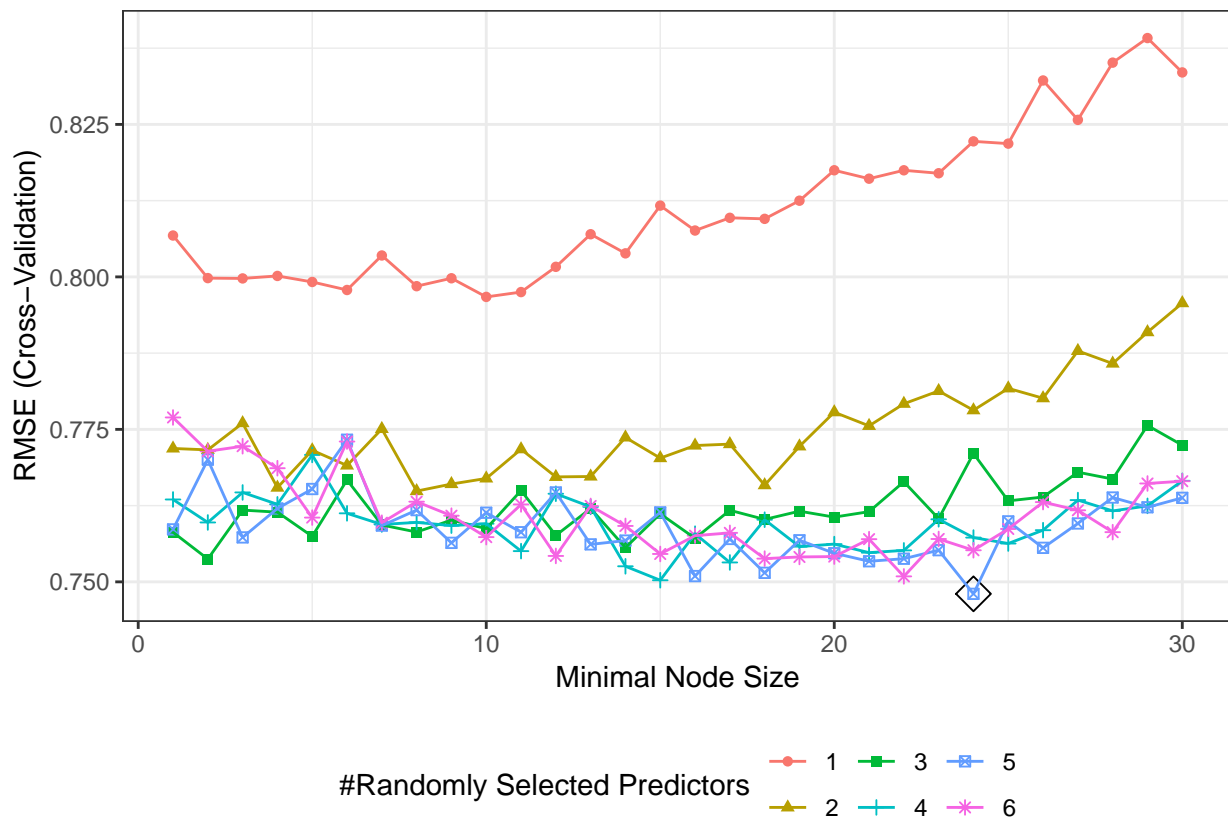
d) Random Forests

```
set.seed(1)

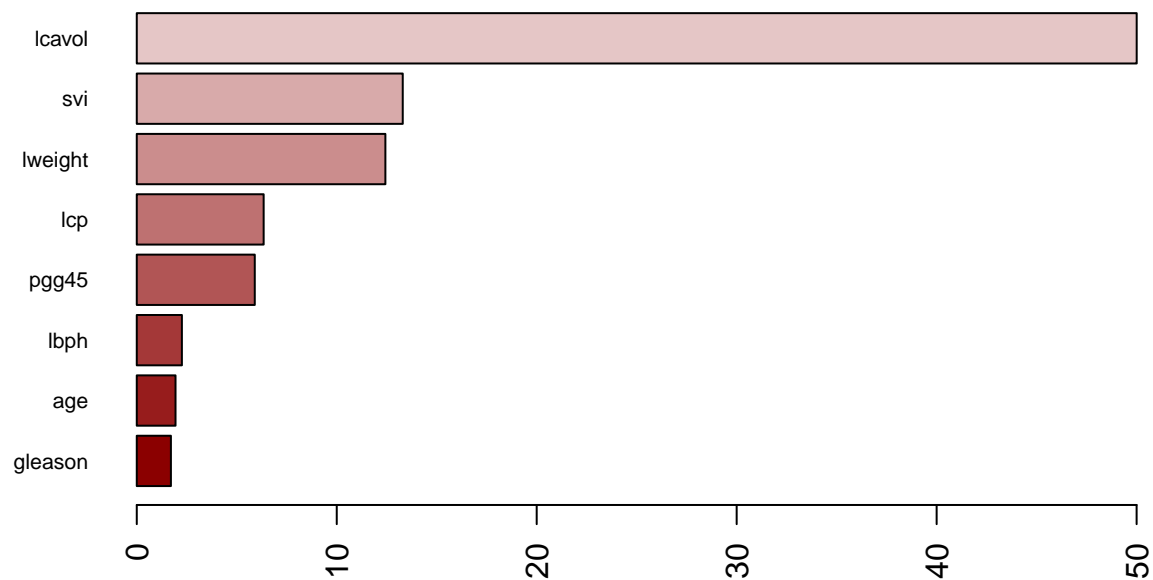
rf.grid <- expand.grid(mtry = 1:6,
                      splitrule = "variance",
                      min.node.size = 1:30
                      )

rf_fit = train(lpsa~.,
               data = Prostate,
               method = "ranger",
               tuneGrid = rf.grid,
               importance = "impurity",
               trControl = ctrl)

ggplot(rf_fit, highlight = T)
```



```
barplot(sort(ranger::importance(rf_fit$finalModel),
  decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(
    colors = c("darkred", "white", "darkblue"))(19))
```



The lcavol is the most importance variable in this random forests model.

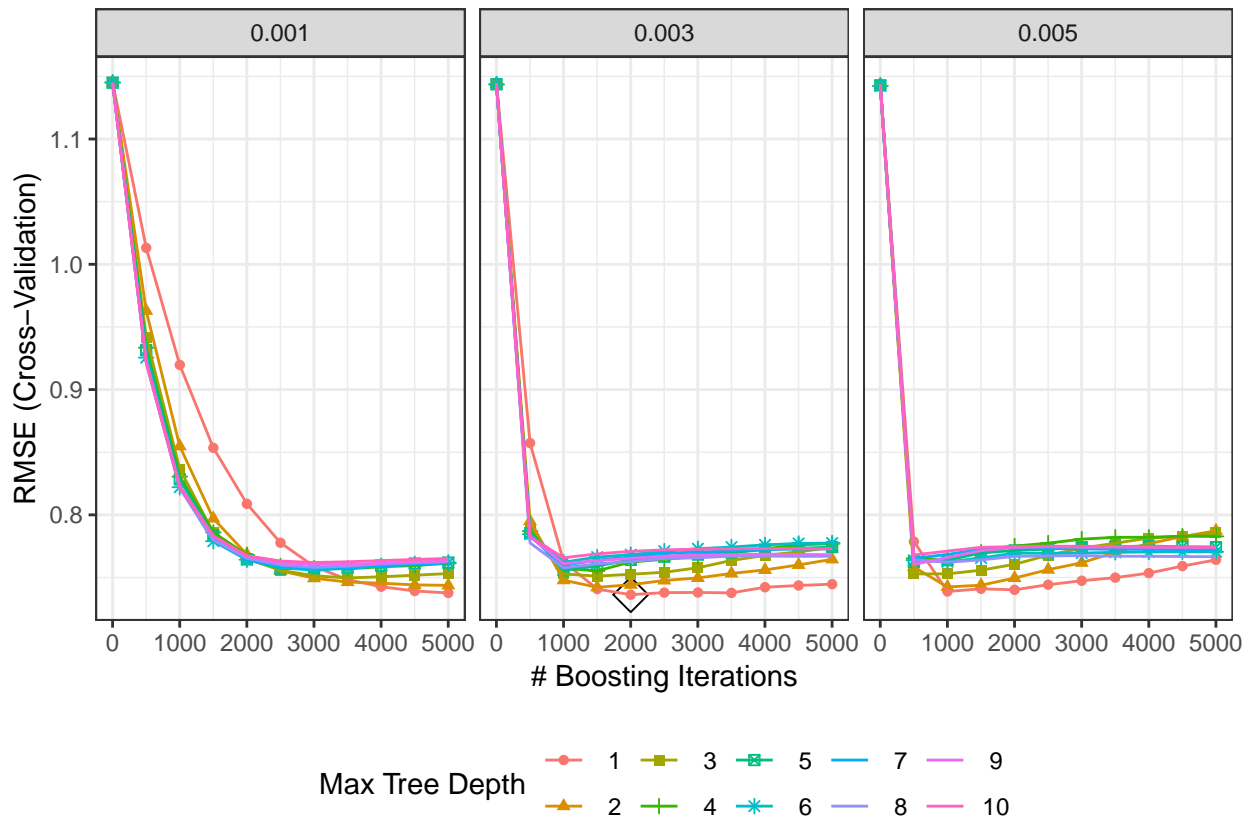
Importance: lcavol > svi > lweight > lcp > pgg45 > lbph > age > gleason

e) Boosting

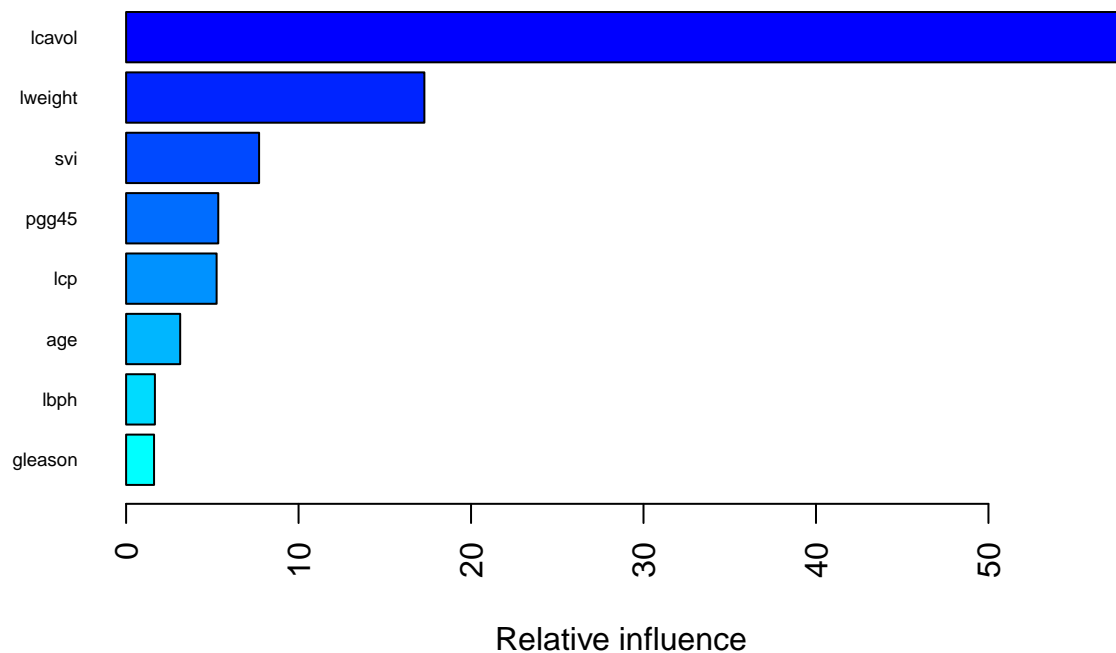
```
set.seed(1)
gbm.grid = expand.grid(
  n.trees = seq(1,5001, by = 500),
  interaction.depth = 1:10,
  shrinkage = c(0.001, 0.003, 0.005),
  n.minobsinnode = 1
)
```

```
gbm_fit = train(lpsa~.,
  data = Prostate,
  method = "gbm",
  tuneGrid = gbm.grid,
  trControl = ctrl,
  verbose = FALSE)
```

```
ggplot(gbm_fit, highlight = T)
```



```
summary(gbm_fit$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```

```
##      var  rel.inf
## lcavol  lcavol 57.977658
## lweight lweight 17.294196
## svi     svi    7.711160
## pgg45   pgg45  5.343532
## lcp     lcp    5.247288
## age     age    3.135510
## lbph    lbph   1.669569
## gleason gleason 1.621087
```

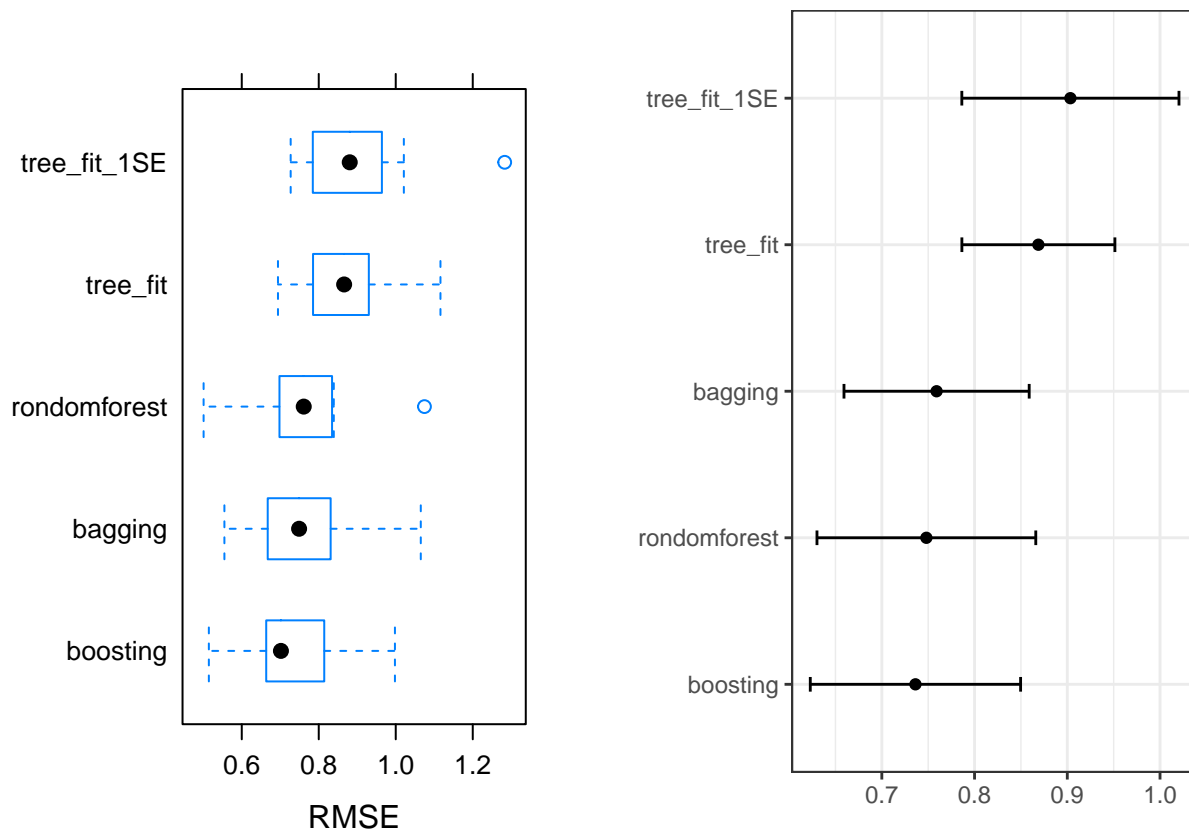
The lcavol is the most importance variable in this boosting model.

Importance: lcavol > lweight > svi > lcp > pgg45 > age > lbph > gleason

d) Compare Models

```
resamp <- resamples(list(tree_fit = tree_fit,
                        tree_fit_1SE = tree_fit_2,
                        bagging = bag_fit,
                        rondonforest = rf_fit,
                        boosting = gbm_fit
                        ))
a = bwplot(resamp, metric = "RMSE")
b = ggplot(resamp, metric = "RMSE")

gridExtra::grid.arrange(a,b,ncol = 2,nrow = 1)
```



From the boxplots of RMSE in the cross-validation, we can see that ensemble methods (bagging, random forest and boosting) have a better performance in the cross-validation than the simple decision tree model. Comparing means of RMSE of different models in the cross-validation, we choose the boosting model to predict PSA level as it has the lowest mean.

Problem 2

a) Decision Tree

```
data(OJ)

set.seed(1)

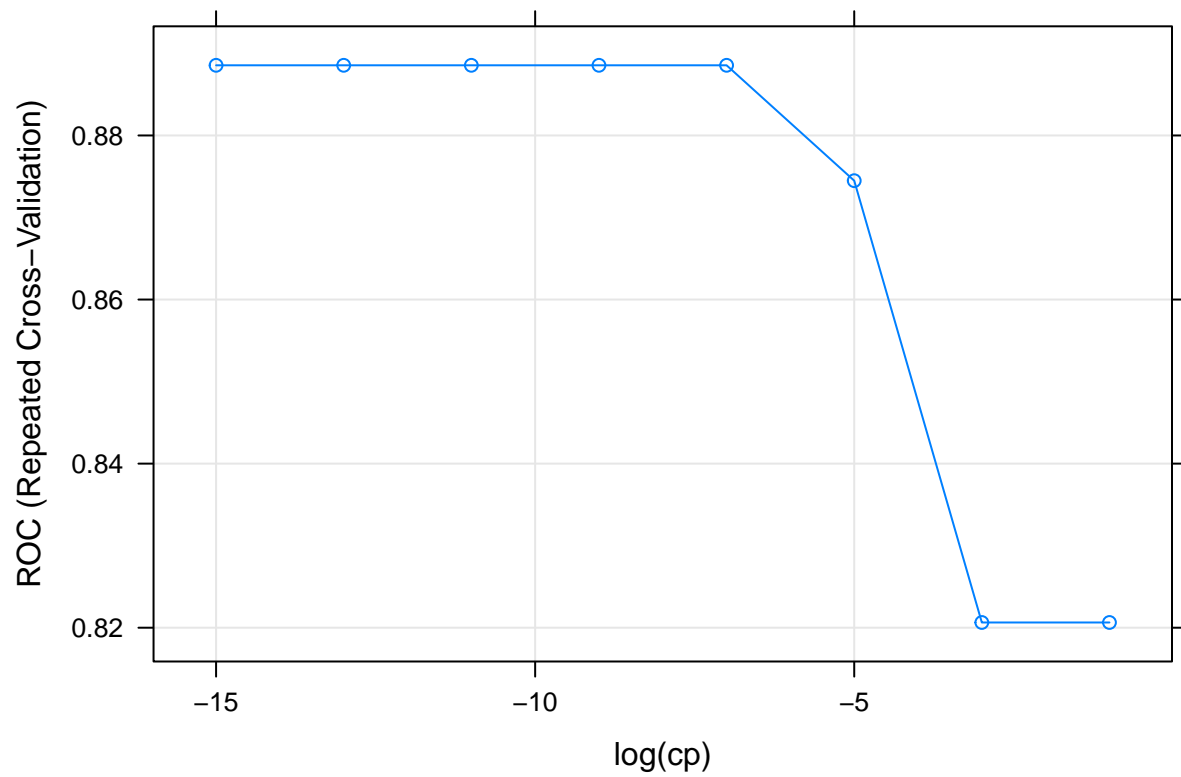
train_ind <- sample(seq_len(nrow(OJ)), size = 800)

training <- OJ[train_ind, ]
test <- OJ[-train_ind, ]

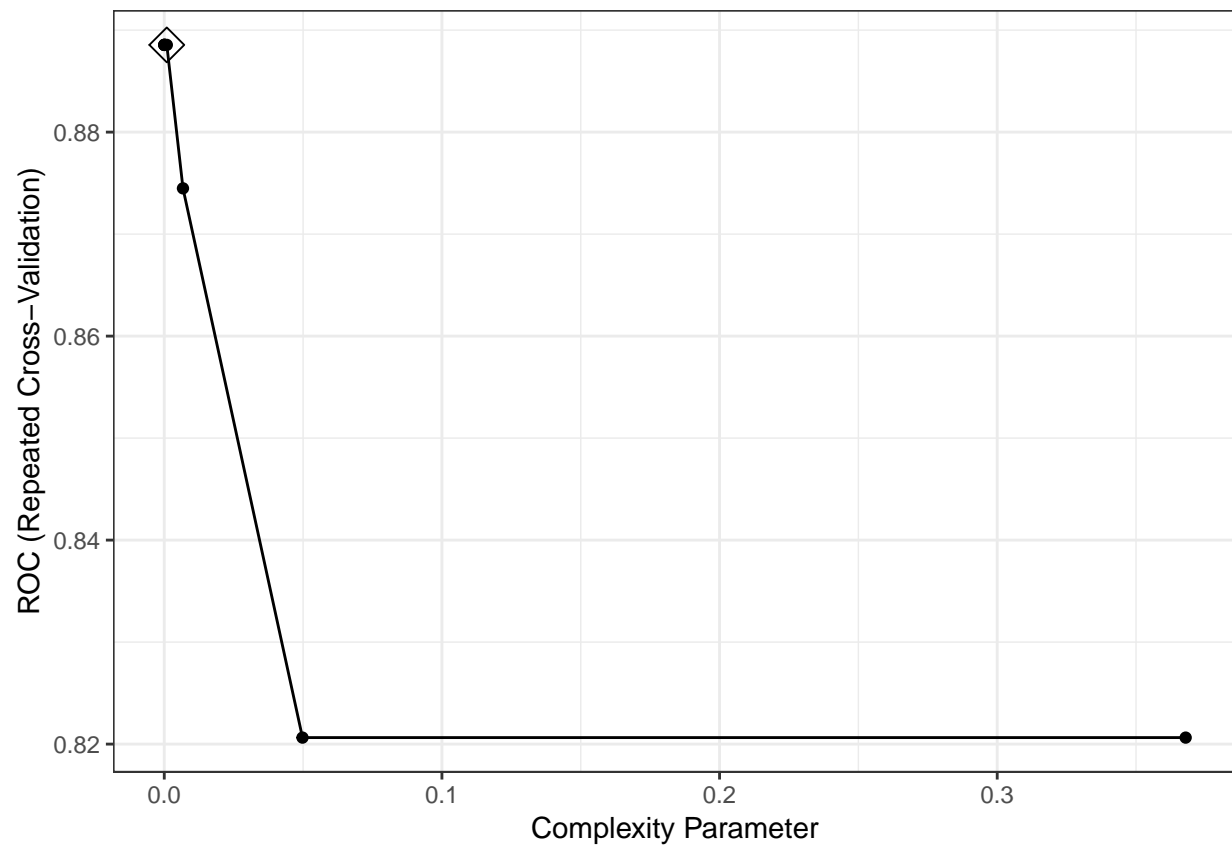
set.seed(1)
ctrl = trainControl(method = "repeatedcv",
                    summaryFunction = twoClassSummary,
                    classProbs = TRUE)

tree_fit_c = train(Purchase~.,
                  data = training,
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-15,0, by = 2))),
                  trControl = ctrl,
```

```
metric = "ROC"  
)  
  
plot(tree_fit_c, xTrans = function(x)log(x), xlab = "log(cp)")
```



```
ggplot(tree_fit_c, highlight = T)
```



```
tree_fit_c$bestTune
```

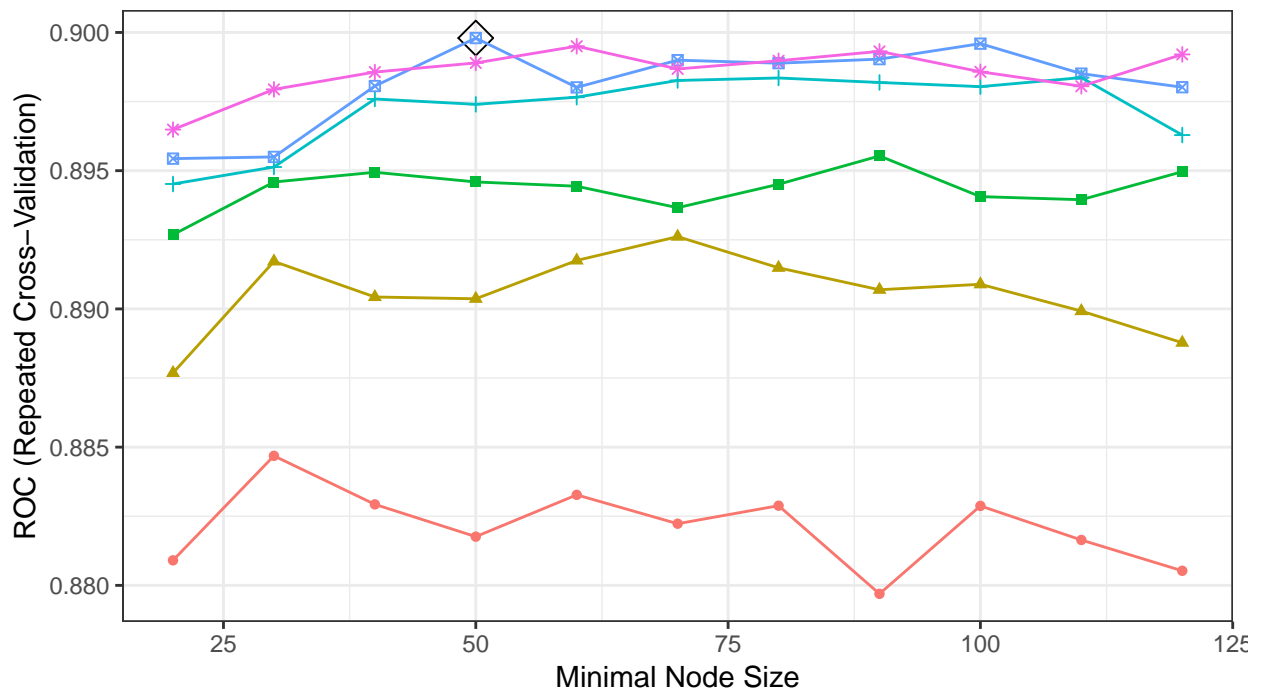
```
##          cp
## 5 0.000911882
```

```
tree_fit_c$finalModel$cptable
```

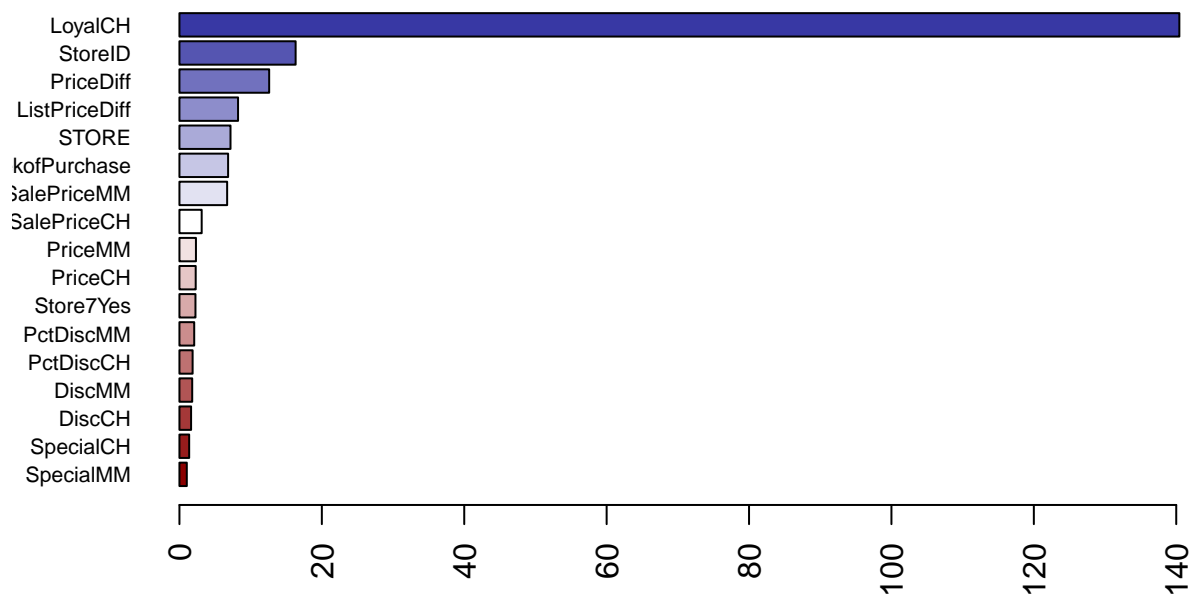
```
##          CP nsplit rel error
## 1 0.522875817    0 1.0000000
## 2 0.016339869    1 0.4771242
## 3 0.009803922    4 0.4281046
## 4 0.006535948    8 0.3888889
## 5 0.003267974   14 0.3496732
## 6 0.001633987   15 0.3464052
## 7 0.001089325   19 0.3398693
## 8 0.000000000   22 0.3366013
```

The plot of the final tree

```
rpart.plot(tree_fit_c$finalModel)
```

```
barplot(sort(ranger::importance(rf_fit_c$finalModel),
  decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(
    colors = c("darkred", "white", "darkblue"))(19))
```



```
rf.pred = predict(rf_fit_c, newdata = test, type = "raw")
1 - sum(rf.pred == test$Purchase) / length(test$Purchase)
```

```
## [1] 0.1703704
```

The test classification error rate is 17.04% for this tree model.

The loyalCH is the most importance variable in this boosting model.

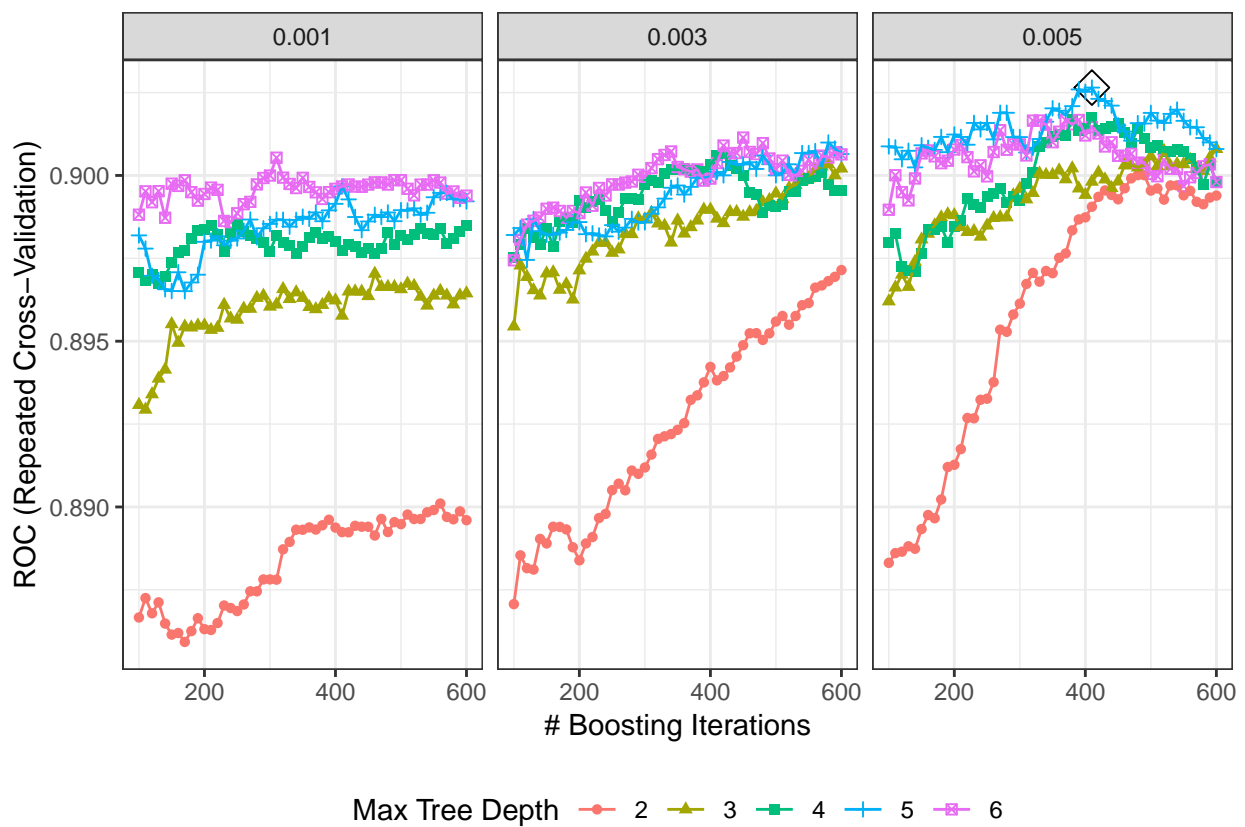
The top 5 most important variables: LoyalCH > StoreID > PriceDiff > ListPriceDiff > STORE

c) Boosting

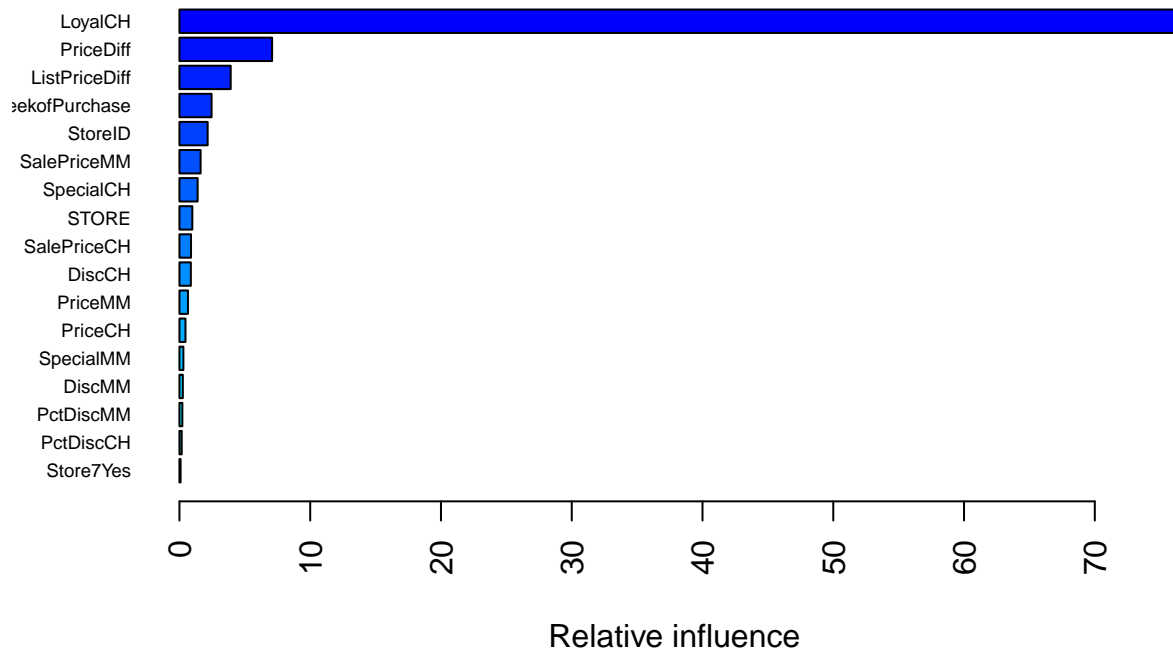
```
set.seed(1)
gbm.grid = expand.grid(n.trees = seq(100, 600, by = 10),
                      interaction.depth = 2:6,
                      shrinkage = c(0.001, 0.003, 0.005),
                      n.minobsinnode = 1)

gbm_fit_c = train(Purchase~.,
                  data = training,
                  method = "gbm",
                  trControl = ctrl,
                  distribution = "bernoulli",
                  metric = "ROC",
                  tuneGrid = gbm.grid,
                  verbose = F
                  )

ggplot(gbm_fit_c, highlight = T)
```



```
summary(gbm_fit_c$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```



```
##           var      rel.inf
## LoyalCH      LoyalCH 76.46936133
## PriceDiff    PriceDiff 7.08748777
## ListPriceDiff ListPriceDiff 3.91872613
## WeekofPurchase WeekofPurchase 2.45739050
## StoreID      StoreID 2.15767257
## SalePriceMM   SalePriceMM 1.61919790
## SpecialCH     SpecialCH 1.39132422
## STORE        STORE 0.98430038
## SalePriceCH   SalePriceCH 0.88406034
## DiscCH        DiscCH 0.87234595
## PriceMM       PriceMM 0.65350459
## PriceCH       PriceCH 0.46000519
## SpecialMM     SpecialMM 0.30142017
## DiscMM        DiscMM 0.25856578
## PctDiscMM     PctDiscMM 0.22307879
## PctDiscCH     PctDiscCH 0.17375185
## Store7Yes     Store7Yes 0.08780651
```

```
gbm.pred = predict(gbm_fit_c, newdata = test, type = "raw")
```

```
1 - sum(gbm.pred == test$Purchase) / length(test$Purchase)
```

```
## [1] 0.1814815
```

The test classification error rate is 18.15% for this boosting model.

The loyalCH is the most importance variable in this boosting model. The top 5 most important variables:
LoyalCH > StoreID > PriceDiff > ListPriceDiff > STORE